

B-Tree and B+Tree

Weiss Book Chapter 4.7

Byoungyoung Lee

<https://compsec.snu.ac.kr>

byoungyoung@snu.ac.kr

Outline


Memory issues in designing data structures

B-Tree

B+Tree


Memory Considerations

- When we discuss data structures, we never specifically mention where the data would be stored.
- In fact, memory hierarchy suggests that the design of data structures should be well aware of it



Memory Considerations

- Memory things to think about when designing data structures
 - Access speed
 - Cost per memory size
 - The unit size of access
 - Stream access vs. Random access



Tree for very large datasets

- Suppose you got very many pieces of information
 - e.g., $n = 2^{30}$
 - Suppose each piece has 1KB data
 - Examples
 - Student records, where each piece holds each student's report
 - Sales history, where each piece holds sale records per item
- If you design the tree to store such data
 - The number of nodes: 2^{30}
 - Each node will occupy at least 1KB
 - Total? At least 1TB

```
1 class sales_node {  
2     ID item_id;  
3     RECORDS records; // 1KB  
4     sales_node *p_left_tree;  
5     sales_node *p_right_tree  
6     // ...  
7     // more  
8     // ...  
9 }
```

Issues: Tree for very large datasets


- Two performance issues of a “very large” tree
 - **Problem #1: Too many memory access**
 - Memory access is slow
 - But the number of memory access would be the height of the tree
 - Because you only load a single node per memory access
 - **Problem #2: Too big to store in the fast memory**
 - Can you store 1TB in the fast memory?
 - No, your main memory is (very likely) smaller than 1TB
 - So you will need to store the tree in the slow memory (i.e., a disk)

Ideas

- **Solution #1: Multiple keys per node**
 - Load multiple keys at once
 - Reduce the height of the tree
 - Trees with this feature
 - M-Way Search Trees, B-Tree, B+Tree
- **Solution #2: No data in the internal nodes**
 - Leaf nodes hold both keys/data, and internal nodes only hold keys
 - Entire “internal nodes” can be stored in the fast memory
 - Trees with this feature
 - B+Tree


M-Way Search Trees

M-Way Search Trees: A search tree with maximum branching factor M



B-Trees

- Each node has keys up to $M-1$ keys
- Order property
 - Subtree between two keys x and y contain leaves with values v such that $x < v < y$




B-Tree Structure Property

- Root (special case)
 - Has between 2 and M children (or root could be a leaf node)
- Internal nodes
 - Store up to $M - 1$ keys
 - Have between $\lceil M/2 \rceil$ and M children
- Leaf nodes
 - Store between $\lceil (M - 1)/2 \rceil$ and $M - 1$ sorted keys
 - All at the same depth
- Note: Our B-Tree assignment has a slight different structural property, but the key ideas are the same

B-Tree: Example

- B-Tree with **M = 3**



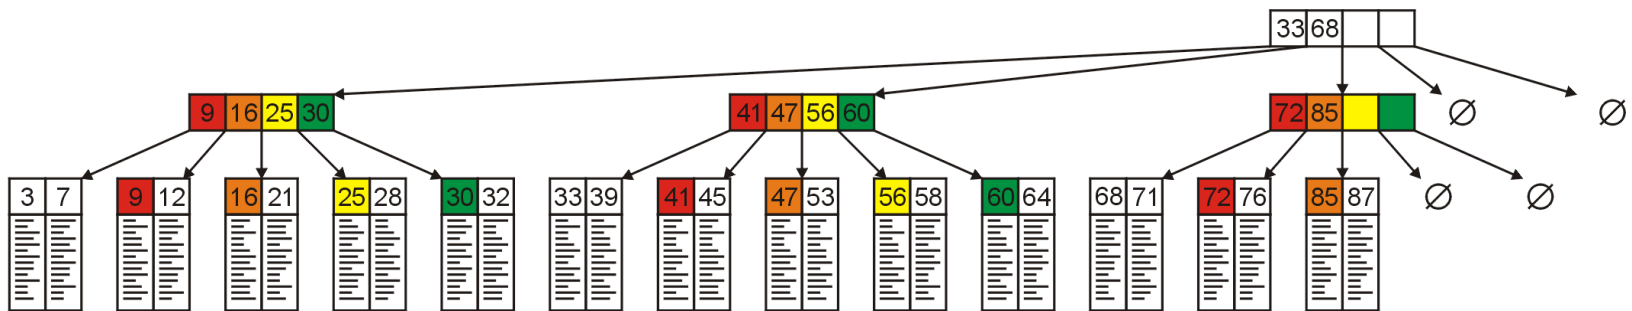
B+Trees

- Internal nodes have no data
 - Only leaf nodes have data
- Each internal node still has (up to) $M - 1$ keys
- Order property
 - Subtree between two keys x and y contain leaves with values v such that $x \leq v < y$
 - Note the symbol, ' \leq '
- Leaf nodes have up to L sorted keys

B+Tree Structure Property

- Root (special case)
 - Has between 2 and M children (or root could be a leaf node)
- Internal nodes
 - Store up to $M - 1$ keys
 - Have between $\lceil M/2 \rceil$ and M children
- Leaf nodes
 - Where data is stored
 - All at the same depth
 - Contain between $\lceil L/2 \rceil$ and L data items

B+Tree: Example



Disk Friendliness of B+Tree

- Many keys stored in a node
 - All brought to memory/cache in one disk access
- Internal nodes contain only keys
 - Only leaf nodes contain actual data
 - Much of tree structure can be loaded into memory irrespective of data object size
 - Data actually resides in disk

Performance Comparison: B+Tree vs. AVL Tree

- Suppose again you have $n = 2^{30}$ items
 - AVL Tree
 - Height: 43
 - B+Tree where $M=256$, $L=256$
 - Height: 4.3
- If you consider other factors, things are getting more interesting
 - The size of each item
 - The size of Cache
 - The size of DRAM
 - ...
- We never talked about the costs to balance the tree though

Maintain the Balance of B+ Tree

- How to make B+ Tree balanced?
 - Insertion idea (bottom-up approach)
 - Step 1: Insert an item to the leaf
 - Step 2: If the node overflows, 1) split the node and 2) add the key to the parent
 - Step 3: If the parent overflows, go back to step 2
 - You may need to increase the height
 - Deletion idea (bottom-up approach)
 - Step 1: Remove an item from the leaf
 - Step 2: If the node underflows, 1) adopt from (or merge with) the neighbor and 2) update the parent
 - Step 3: If the parent underflows, go back to step 2
 - You may need to decrease the height

Applications

- Databases
 - Index structure for MySQL
 - A hash table can be a better option (will cover later)
- File systems
 - Apple's HFS+, Microsoft NTFS, Linux's EXT4 and btrfs
- Real-world challenges in designing and implementing trees
 - Parallel access: Multi-core processors are everywhere
 - Distributed storage: Too large to store in a single computing node
 - You will learn more from advanced courses: operating systems, computer architecture, database, etc.

References

Wikipedia, http://en.wikipedia.org/wiki/B+_tree

- [1] Cormen, Leiserson, and Rivest, *Introduction to Algorithms*, McGraw Hill, 1990, Ch. 19, p.381-99.
- [2] Weiss, *Data Structures and Algorithm Analysis in C++*, 3rd Ed., Addison Wesley, §4.7, p.159-64.
- [3] Lecture slides by Brian Curless -
<https://courses.cs.washington.edu/courses/cse326/08sp/lectures/markup/11-b-trees-markup.pdf>