

An Introduction to Hash Tables

Weiss Book Chapter 5

Byoungyoung Lee

<https://compsec.snu.ac.kr>

byoungyoung@snu.ac.kr

Outline

Discuss storing unrelated/unordered data

- IP addresses and domain names

Consider conversions between these two forms

Introduce the idea of hashing:

- Reducing $\mathbf{O}(\ln(n))$ operations to $\mathbf{O}(1)$

Consider some of the weaknesses

Problem: IP Addresses

Examples:

You want to map an IP address to a corresponding domain name

A 32-bit IP address are often written as four byte values from 0 to 255

- Consider IP Address
 - $10010011\ 00101110\ 01111001\ 00010110_2$
 - This can be written as 147.46.121.22
- Suppose its domain name is
 - `compsec.snu.ac.kr`
- We use domain names because IP addresses are not human readable

Example: IP Addresses

Given an IP address, if we wanted to quickly find any associated domain name, we could create an array of size 2^{32} (4294967296) of strings:

```
int const MAX_IP_ADDRESSES = 4294967296;  
string domain_name[MAX_IP_ADDRESSES];
```

For example, 147.46.121.22 can be translated

– As $147 * 256^3 + 46 * 256^2 + 121 * 256 + 22 = 2466293526$,

```
domain_name[2466293526] = "compsec.snu.ac.kr";
```

Can we use much less memory than this?

Goals and Requirements

Our goal:

Store data so that all operations are $\Theta(1)$ time

Requirement:

The memory requirement should be $\Theta(n)$

Q. Can we achieve this goal with data structures we covered before?

- Lists, stack, queue, trees, ...

Goals and Requirements

In general, we would like to:

- Create an array of size M
- Store each of n objects in one of the M bins
- Have some means of determining the bin in which an object is stored

Idea: Grade Table Example

Let's try a simpler problem

- How do I store your examination grades so that I can access your grades in $\Theta(1)$ time?

Observation: SNU ID is an 9-digit number

- I can't create an array of size 10^9
- I can create an array of size 1000 though
- How could you convert an 9-digit number into a 3-digit number?
- First three digits might cause a problem
 - almost all students start with 2017, 2018, 2019, ...
- The last four digits, however, are (somehow) random

Therefore, I could store the examination grade for SNU ID 201901011

```
grade[011] = 99;
```

Idea: Grade Table Example

Consequently, I have a function, mapping a student ID to a 3-digit number

- I can store the information in that location
- Storing it, accessing it, and erasing may take $\Theta(1)$
- Problem: two or more students may map to the same number:
 - Vayne has ID 200703456
 - Teemo has ID 200301456
 - Both would map to 456
- ➔ So called “collision”

⋮	⋮
454	
455	
456	86
457	
458	
459	
460	
461	
462	
463	79
464	
465	
⋮	⋮

Probability of Collision

Question:


- What is the likelihood that in a class of size 100 that **no two students will have the same last three digits**?
- Not very high: i.e., a probability of having “no collision”

$$1 \cdot \frac{999}{1000} \cdot \frac{998}{1000} \cdot \frac{997}{1000} \cdot \dots \cdot \frac{901}{1000} \approx 0.005959$$

- Probability of having collision(s): $1 - 0.005959 = 0.994041$
- Implication: If you insert 100 students to the table, there will be at least one collision at the probability of more than 99.4%
 - So highly likely there will be **a collision** if only using the last three digits

Check more:

https://en.wikipedia.org/wiki/Birthday_problem



The hashing problem

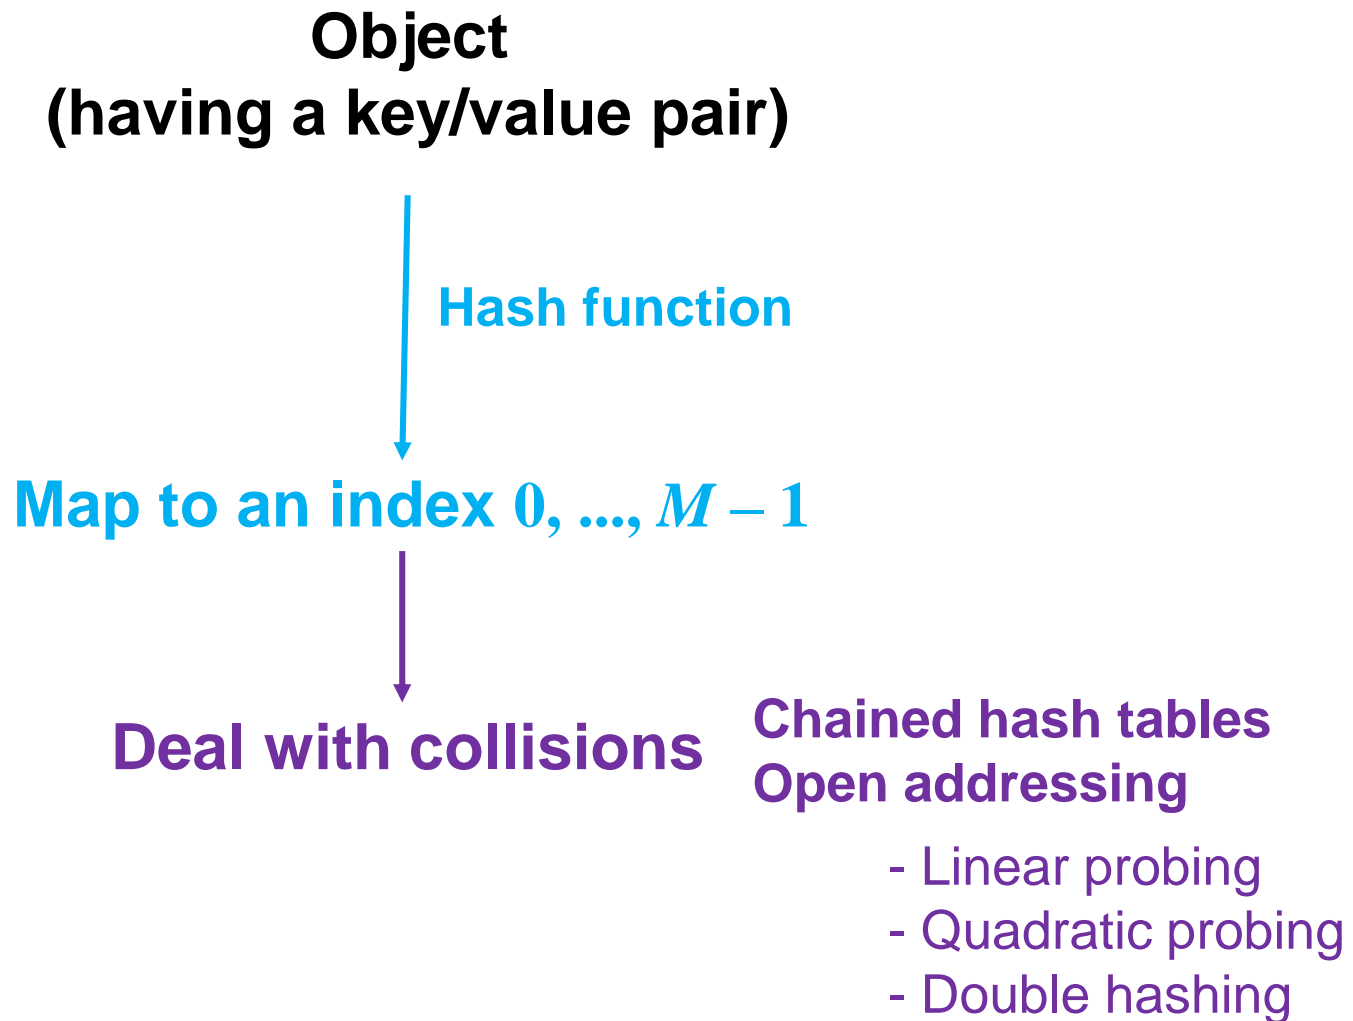
The process of mapping an object or a number onto an integer in a given range is called *hashing*

Problem: multiple objects may hash to the same value

- Such an event is termed *a collision*

Hash tables use a hash function together with a mechanism for dealing with collisions

The hash process



Definition: Hash Function

What is a hash of an object?


From Oxford Languages:

a dish of cooked meat cut into small pieces and cooked again, usually with potatoes.

Our definition: map an arbitrary-size input to an fixed-size value (e.g., **M**)
 $0, 1, 2, \dots, M - 1$



The hash process



Ideal properties of a hash function

A hash function is a function mapping an input key to a certain range (say 0 to M)

Necessary properties of such a hash function h are:

- 1a. Computation should be **fast**, ideally $\Theta(1)$
- 1b. The hash value must be **deterministic**
 - It must always return the same output: If $x = y \Rightarrow h(x) = h(y)$
- 1c. If two objects are randomly chosen, there should be only $1/M$ chance that they have the same hash value
(i.e., **uniform distribution**)

Types of hash functions

Hash functions for different types of input keys

- General class object
- Integer
- String

Hash Function for Class Object #1

#1. The easiest solution is to give each object **a unique number**

```
class Product {  
    private:  
        unsigned int hash_value;  
        static unsigned int hash_count;  
    public:  
        Product();  
        unsigned int hash() const;  
};  
  
unsigned int Product::hash_count = 0;
```

```
Product::Product() {  
    hash_value = hash_count;  
    ++hash_count;  
}  
  
unsigned int Product::hash() const {  
    return hash_value;  
}
```


Hash Function for Class Object #2

#2. If we only need the hash value while the object exists in memory, you may use the **address**:

```
unsigned int Product::hash() const {  
    return reinterpret_cast<unsigned int>( this );  
}
```

Check more: https://en.cppreference.com/w/cpp/language/reinterpret_cast

Hash Function for Integer


- Knuth's Multiplicative Method

$$\text{hash}(i) = i * 2654435761 \bmod 2^{32}$$

- 2654435761 is the golden ratio of 2^{32}
- 2654435761 and 2^{32} have no common factors
 - So the multiplication produces a complete mapping of the key to hash result with no overlap
 - Having common factor “n” would only map to $1/n$ possible hashes
- Issue: This preserves the divisibility. So for example, if your keys were even, their hashes are always even too.
- Proof the claim:

Consider $h(x) = ax \bmod m$.

If a and m are coprime, $h(x_1) \neq h(x_2)$ if $x_1 < x_2 < m$.



Hash Function for Integer

- Robert Jenkin's 32-bit integer hash function

```
uint32_t hash( uint32_t a)
{
    a = (a+0x7ed55d16) + (a<<12);
    a = (a^0xc761c23c) ^ (a>>19);
    a = (a+0x165667b1) + (a<<5);
    a = (a+0xd3a2646c) ^ (a<<9);
    a = (a+0xfd7046c5) + (a<<3);
    a = (a^0xb55a4f09) ^ (a>>16);
    return a;
}
```

Hash Function for Integer

- It's difficult to tell if your hash function is good or bad
- It depends on the distribution of input keys
- What should be the goodness measure of hash function?
 - The hash function must be efficiently computed
 - The key mapping is closed to the uniform distribution

Hash Function for String

Two strings are equal if all the characters are equal and in the identical order

A string is simply an array of bytes:

- Each byte stores a value from 0 to 255

Any hash function must be a function of these bytes

Hash Function for String

We could, for example, just add the characters:

```
unsigned int hash( const string &str ) {  
    unsigned int hash_value = 0;  
  
    for ( int k = 0; k < str.length(); ++k ) {  
        hash_value += str[k];  
    }  
  
    return hash_value;  
}
```

Q. What's the problem of this hash function?

Hash Function for String

Let the individual characters represent the coefficients of a polynomial in x :


$$p(x) = c_0 x^{n-1} + c_1 x^{n-2} + \dots + c_{n-3} x^2 + c_{n-2} x + c_{n-1}$$

Use Horner's rule to evaluate this polynomial at a prime number, e.g., $x = 12347$:

```
unsigned int hash( string const &str ) {  
    unsigned int hash_value = 0;  
  
    for ( int k = 0; k < str.length(); ++k ) {  
        hash_value = 12347*hash_value + str[k];  
    }  
    return hash_value;  
}
```

Note: Hash functions here != Cryptographic Hash

- All the hash functions discussed here are not cryptographic hash functions
 - MD5, SHA-1, SHA-512
 - https://en.wikipedia.org/wiki/Cryptographic_hash_function



<https://en.bitcoinwiki.org/wiki/SHA-256>

- Cryptographic hash functions have following security properties
 - Pre-image resistance
 - Given h , it is difficult to find m such that $h = \text{hash}(m)$
 - Second pre-image resistance
 - Given m_1 , it is difficult to find m_2 such that $\text{hash}(m_1) = \text{hash}(m_2)$
 - Collision resistance
 - Difficult to find any m_1 and m_2 such that $\text{hash}(m_1) = \text{hash}(m_2)$
 - Be careful on the definition of being “difficult”



<https://blog.bankofhodlers.com/the-value-of-proof-of-work/>

Mapping down to 0, ..., $M-1$

By far, we computed 32-bit hash values for different input keys

- Class object
- Integer
- String

Practically, we will require a hash value on the range 0, ..., $M-1$:

- The modulus operator %

```
unsigned int hash_M( unsigned int n, unsigned int M ) {  
    return n % M;  
}
```

```
unsigned int hash_M( unsigned int n, unsigned int m ) {  
    return n & ((1 << m) - 1); // M = 2^m  
}
```

Summary

Hashing

- Discuss storing unordered data
- Discuss IP addresses and domain names
- Discussed the issues with collisions

Hash Function

- Ideal properties
- Hash functions for
 - Integer
 - String
 - Class object
- Map a 32-bit integer onto a smaller range $0, 1, \dots, M - 1$

References

Wikipedia, http://en.wikipedia.org/wiki/Hash_table

- [1] Cormen, Leiserson, and Rivest, *Introduction to Algorithms*, McGraw Hill, 1990.
- [2] Weiss, *Data Structures and Algorithm Analysis in C++*, 3rd Ed., Addison Wesley.