

Abstract Priority Queues

Textbook: Weiss Chapter 6.1, 6.2

Byoungyoung Lee

<https://compsec.snu.ac.kr>

byoungyoung@snu.ac.kr

Outline

This topic will:

- Review queues
- Discuss the concept of priority and priority queues
- Look at two simple implementations:
 - Arrays of queues
 - AVL trees
- Introduce heaps, an alternative tree structure which has better run-time characteristics

Background

We have discussed Abstract Lists

- Arrays, linked lists

We saw three cases which restricted the operations:

- Stacks, queues, dequeues

Then, we studied search trees: Abstract Sorted Lists

- Run times were generally $\Theta(\ln(n))$

We will now look :

- Priority queues
- Restriction on Abstracted Sorted Lists

Definition

With queues

- The order may be summarized by *first in, first out*

If each object is associated with a priority, we may wish to pop that object which has highest priority

With each pushed object, we will associate a nonnegative integer (0, 1, 2, ...) where:


- The value 0 has the *highest* priority, and
- The higher the number, the lower the priority

Operations


The top of a priority queue is the object with highest priority



Popping from a priority queue removes the current highest priority object:



Push places a new object into the appropriate place



Process Priority in Linux

This is the scheme used by Linux, e.g.,

```
% nice -15 ./a.out
```

sets the priority of the execution of a.out as 15

The kernel will schedule processes according to the priority

```
$ man nice
```

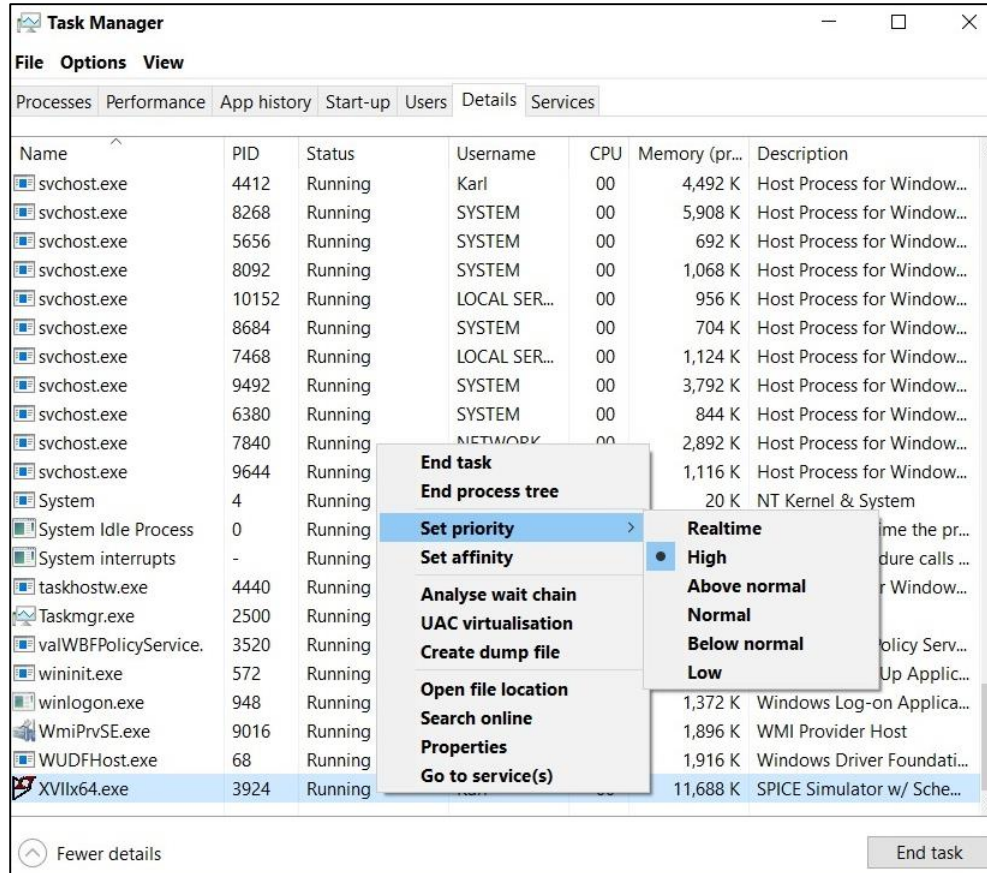
```
NICE(1)                                User Co
ands                                     NICE(1)

NAME
    nice - run a program with modified scheduling priority

SYNOPSIS
    nice [OPTION] [COMMAND [ARG]...]

DESCRIPTION
    Run COMMAND with an adjusted niceness, which affects process scheduling. With no COMMAND, print the current niceness.
    Niceness values range from -20 (most favorable to the process) to 19 (least favorable to the process).
```

Process Priority in Windows



Implementations

Our goal is to make the run time of each operation as close to $\Theta(1)$ as possible

We will look at two naïve implementations using data structures we already know:

- Multiple queues—one for each priority
- An AVL tree

Multiple Queues

Assume there is a fixed number of priorities, say M

- Create an array of M queues
- Push a new object onto the queue corresponding to the priority
- Top and pop find the first empty queue with highest priority

Multiple Queues

The run times are reasonable:

- Push is $\Theta(1)$
- Top and pop are both $O(M)$

Unfortunately:

- It restricts the range of priorities

AVL Trees

We could simply insert the objects into an AVL tree where the order is given by the stated priority:

- Insertion is $\Theta(\ln(n))$
- Top is $\Theta(\ln(n))$
- Remove is $\Theta(\ln(n))$

There is significant overhead for maintaining both the tree and the corresponding balance


Better Idea: Heaps

Can we do better?

- That is, can we reduce some (or all) of the operations down to $\Theta(1)$?

The next topic defines a *heap*

- A tree with the top object at the root
- We will look at binary heaps
- Numerous other heaps exists:
 - d -ary heaps
 - Leftist heaps
 - Skew heaps
 - Binomial heaps
 - Fibonacci heaps
 - Bi-parental heaps



Summary

This topic:

- Introduced priority queues
- Considered two obvious implementations:
 - Arrays of queues
 - AVL trees
- Discussed the run times and claimed that a variation of a tree, a heap, can do better

References

Cormen, Leiserson, Rivest and Stein,
Introduction to Algorithms, The MIT Press, 2001, §6.5, pp.138-44.

Mark A. Weiss,
Data Structures and Algorithm Analysis in C++, 3rd Ed., Addison Wesley, 2006, Ch.6, p.213.

Joh Kleinberg and Eva Tardos,
Algorithm Design, Pearson, 2006, §2.5, pp.57-65.

Elliot B. Koffman and Paul A.T. Wolfgang,
Objects, Abstractions, Data Structures and Design using C++, Wiley, 2006, §8.5, pp.489-96