

Heap sort

Textbook: Weiss Chapter 7.5, Chapter 7.8

Byoungyoung Lee

<https://compsec.snu.ac.kr>

byoungyoung@snu.ac.kr

Outline

This topic covers the simplest $\Theta(n \ln(n))$ sorting algorithm: *heap sort*

We will:

- define the strategy
- analyze the run time
- convert an unsorted list into a heap
- cover some examples

Bonus: in-place sorting

Heap Sort

- Let's use min-heap to implement sorting
 - Inserting n objects into a min-heap first
 - Then popping n objects will pop in sorted order
- Heap sort strategy
 - Given an unsorted list with n objects, place them into a heap, and take them out

Run time Analysis of Heap Sort

Taking an object out of a heap with n items requires $\mathbf{O}(\ln(n))$ time

Therefore, taking n objects out requires

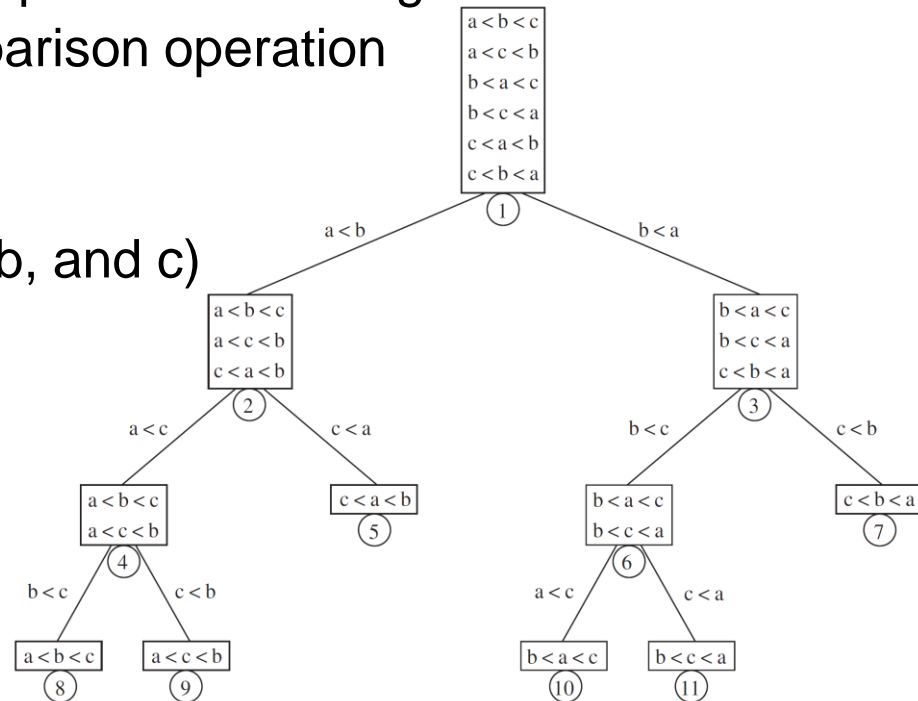
$$\sum_{k=1}^n \ln(k) = \ln\left(\prod_{k=1}^n k\right) = \ln(n!)$$

Q. What is the asymptotic bound of $\ln(n!)$?

Let's step back, and think about the runtime of generic sorting algorithms

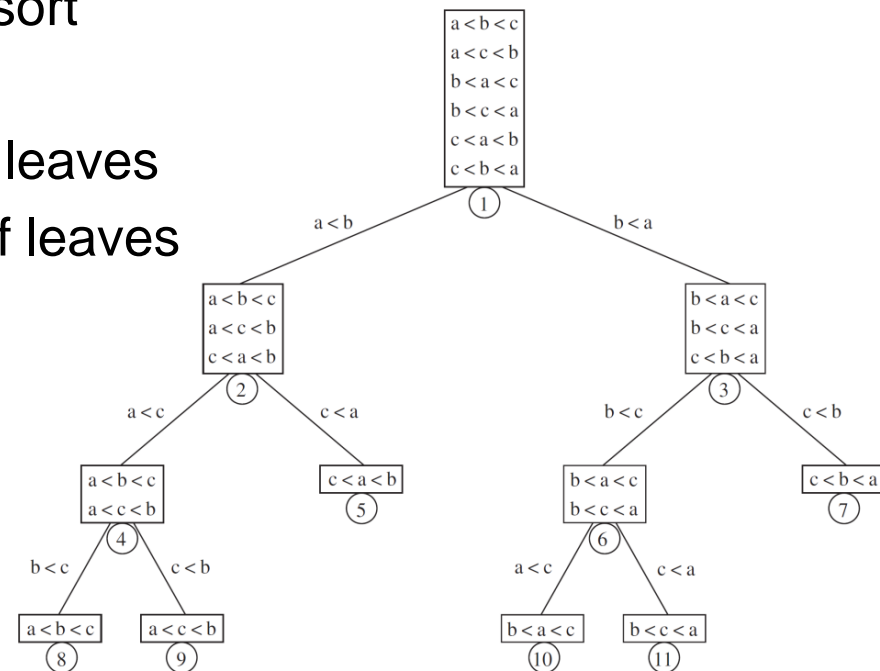
General Lower Bound of Sorting

- **A decision tree**
 - Abstract representation of **any comparison-based sorting algorithm**
 - A decision tree is a binary tree
 - Each node represents a set of possible ordering
 - Each edge represents a comparison operation
 - The right shows the decision tree
 - for three-element sort (i.e., a, b, and c)



General Lower Bound of Sorting

- **A decision tree**
 - Any comparison-based sorting algorithm can be represented with the decision tree
 - The depth of a leaf node
 - The number of comparisons to sort
 - Worst case: the maximum depth of leaves
 - Average case: the average depth of leaves



General Lower Bound of Sorting

- **Lemma (7.1, Weiss p.324)**
 - Let T be a binary tree of depth d . Then T has at most 2^d leaves
- **Proof by induction**
 - **Base case**
 - If $d = 0$, there's at most one leaf, which is true
 - **Inductive step**
 - Suppose T has two subtrees and the depth of T is d
 - T must have two subtrees, where the depth of each subtree is at most $d-1$
 - Assuming the induction hypothesis is true for $d-1$, each subtree has at most 2^{d-1} leaves
 - T has at most 2^d leaves, which proves the lemma.

General Lower Bound of Sorting

- **Lemma (7.2)**
 - A binary tree with L leaves must have depth at least $\lceil \log L \rceil$
 - **Proof**
 - Easy to proof based on Lemma 7.1
- **Theorem (7.6)**
 - Any sorting algorithm that uses only comparisons requires at least $\lceil \log(N!) \rceil$ comparisons in the worst case
 - **Proof:** A decision tree to sort N elements must have $N!$ leaves
 - **Q. why does the decision tree have $N!$ leaves?**

General Lower Bound of Sorting

- **Theorem (7.7)**

- Any sorting algorithm that uses only comparisons requires $[N \log N]$ comparisons

- **Proof**

- $$\begin{aligned} \log(N!) &= \log\{N(N-1)(N-2) \cdots (1)\} \\ &= \log N + \log(N-1) + \log(N-2) + \cdots + \log 1 \\ &\geq \log N + \log(N-1) + \log(N-2) + \cdots + \log\left(\frac{N}{2}\right) \\ &\geq \frac{N}{2} \log\left(\frac{N}{2}\right) \\ &\geq \frac{N}{2} \log N - \frac{N}{2} \\ &= \Omega(N \log N) \end{aligned}$$

Recall: $f(n) = \mathbf{O}(g(n)) \Leftrightarrow g(n) = \mathbf{\Omega}(f(n))$

HeapSort: In-place Implementation

Problem:

- The original heap sort solution was out-of-place sorting
- It requires additional memory, a min-heap of size n
- This requires $\Theta(n)$ memory and is therefore not in place


Is it possible to perform a heap sort in-place?

- using at most $\Theta(1)$ memory (a few extra variables)

In-place Implementation

Instead of implementing a min-heap, consider a max-heap:

- A heap where the maximum element is at the top of the heap and the next to be popped




In-place Heapification

Now, consider this unsorted array:

46	52	28	17	3	63	34	81	70	95
----	----	----	----	---	----	----	----	----	----

This array represents the following complete tree:



This is neither a min-heap, max-heap, or binary search tree

Goal: In-place conversion of this complete tree into a max heap

In-place Heapification: Two strategies


Two strategies:

- **Top-down**

- Assume 46 alone is a max-heap
- Keep inserting the next element into the existing heap
- Similar to the strategy for insertion sort


- **Bottom-up**

- Assume all leaf nodes are already max heaps
- Make corrections so that previous nodes also form max heaps



In-place Heapification: Bottom-up

Let's work bottom-up: each leaf node is a max heap on its own




In-place Heapification: Bottom-up

Each leaf node alone is a trivial max-heap, so we don't do any.

Then we move up one level.


The subtree with 87 as the root is a max-heap



In-place Heapification: Bottom-up


The subtree with 23 is not a max-heap, but swapping it with 55 creates a max-heap

This process is termed **percolating down**




In-place Heapification: Bottom-up

The subtree with 3 as the root is not max-heap, but we can swap 3 and the maximum of its children: 86




In-place Heapification: Bottom-up

Starting with the next higher level, the subtree with root 48 can be turned into a max-heap by swapping 48 and 99




In-place Heapification: Bottom-up

Similarly, swapping 61 and 95 creates a max-heap of the next subtree




In-place Heapification: Bottom-up

As does swapping 35 and 92




In-place Heapification: Bottom-up

The subtree with root 24 may be converted into a max-heap by first swapping 24 and 86 and then swapping 24 and 28




In-place Heapification: Bottom-up

The right-most subtree of the next higher level may be turned into a max-heap by swapping 77 and 99




In-place Heapification: Bottom-up

However, to turn the next subtree into a max-heap requires that 13 be percolated down to a leaf node




In-place Heapification: Bottom-up

The root need only be percolated down by two levels



In-place Heapification: Bottom-up

The final product is a max-heap



Run-time Analysis: Bottom-up Heapify

Considering a perfect tree of height h :

- The maximum number of swaps which a second-lowest level would experience is 1, the next higher level, 2, and so on



Run-time Analysis: Bottom-up Heapify

- At depth k
 - There are at most 2^k nodes
 - Nodes at depth k percolate down at most $h - k$ levels
 - In the worst case, this requires a total of $2^k(h - k)$ swaps

- Writing this sum mathematically, we get:

$$\sum_{k=0}^h 2^k(h - k) = (2^{h+1} - 1) - (h + 1)$$

- Recall that for a perfect tree
 - $n = 2^{h+1} - 1$ and $h + 1 = \lg(n + 1)$

$$\sum_{k=0}^h 2^k(h - k) = n - \lg(n + 1) = \Theta(n)$$

- So bottom-up heapify takes $\Theta(n)$

Runtime Analysis of Heap Sort

1. Heapification runs in $\Theta(n)$
2. Popping n items from a heap of size n , runs in $\Theta(n \ln(n))$ time

So the total algorithm will run in $\Theta(n \ln(n))$ time

Run-time Summary

The following table summarizes the run-times of heap sort

Case	Run Time	Comments
Worst	$\Theta(n \ln(n))$	No worst case
Average	$\Theta(n \ln(n))$	
Best	$\Theta(n \ln(n))$	No best case

Summary

We have seen our first in-place $\Theta(n \ln(n))$ sorting algorithm:

- Convert the unsorted list into a max-heap as complete array
- Pop the top n times and place that object into the vacancy at the end
- It requires $\Theta(1)$ additional memory—it is truly in-place

It is a nice algorithm; however, we will see two other $n \ln(n)$ algorithms;

- Merge sort requires $\Theta(n)$ additional memory
- Quick sort requires $\Theta(\ln(n))$ additional memory

References

Wikipedia, <http://en.wikipedia.org/wiki/Heapsort>

- [1] Donald E. Knuth, *The Art of Computer Programming, Volume 3: Sorting and Searching*, 2nd Ed., Addison Wesley, 1998, §5.2.3, p.144-8.
- [2] Cormen, Leiserson, and Rivest, *Introduction to Algorithms*, McGraw Hill, 1990, Ch. 7, p.140-9.
- [3] Weiss, *Data Structures and Algorithm Analysis in C++*, 3rd Ed., Addison Wesley, §7.5, p.270-4.