

Introduction to Graphs

Textbook: Weiss Chapter 9.1

Byoungyoung Lee

<https://compsec.snu.ac.kr>

byoungyoung@snu.ac.kr

Outline

A graph is an abstract data type for storing adjacency relations

- We start with definitions:
 - Vertices, edges, degree and sub-graphs
- We will describe paths in graphs
 - Simple paths and cycles
- Definition of connectedness
- Weighted graphs
- We will then reinterpret these in terms of directed graphs
- Directed acyclic graphs

Undirected Graph

We will define an Undirected Graph ADT as a collection of **vertices**

$$V = \{v_1, v_2, \dots, v_n\}$$

- The number of vertices is denoted by

$$|V| = n$$

- A collection E of unordered pairs $\{v_i, v_j\}$ is termed **edges** which connect the vertices

Data structures implementing abstract undirected graphs:


- Adjacency matrices
- Adjacency lists

Undirected graphs

Consider this collection of vertices

$$V = \{v_1, v_2, \dots, v_9\}$$

where $|V| = n$




Undirected graphs

Associated with these vertices are $|E| = 5$ edges

$$E = \{\{v_1, v_2\}, \{v_3, v_5\}, \{v_4, v_8\}, \{v_4, v_9\}, \{v_6, v_9\}\}$$

- The pair $\{v_j, v_k\}$ indicates
 - vertex v_j is adjacent to vertex v_k and
 - vertex v_k is adjacent to vertex v_j



Undirected graphs

We will assume in this course that a vertex is not adjacent to itself

- For example, $\{v_1, v_1\}$ will not define an edge

The maximum number of edges in an undirected graph is

$$|E| \leq \binom{|V|}{2} = \frac{|V|(|V|-1)}{2} = O(|V|^2)$$


An undirected graph

Example: given the $|V| = 7$ vertices

$$V = \{A, B, C, D, E, F, G\}$$

and the $|E| = 9$ edges

$$E = \{\{A, B\}, \{A, D\}, \{A, E\}, \{B, C\}, \{B, D\}, \{B, E\}, \{C, E\}, \{C, F\}, \{D, E\}\}$$



Degree


The degree of a vertex is defined as the number of adjacent vertices

$$\text{degree}(A) = \text{degree}(D) = \text{degree}(C) = 3$$

$$\text{degree}(B) = \text{degree}(E) = 4$$

$$\text{degree}(F) = 1$$


$$\text{degree}(G) = 0$$



Those vertices adjacent to a given vertex are its *neighbors*




Sub-graphs

A sub-graph of a graph is a subset of the vertices and a subset of the edges in the original graph



Vertex-induced sub-graphs

A *vertex-induced sub-graph* is a subset of the vertices where the edges are all relevant edges in the original graph



Paths

A path in an undirected graph is **an ordered sequence of vertices**

$(v_0, v_1, v_2, \dots, v_k)$


where $\{v_{j-1}, v_j\}$ is an edge for $j = 1, \dots, k$

- Termed *a path from v_0 to v_k*
- The length of this path is k

Paths

A path of length 4:


(A, B, E, C, F)



Paths

A path of length 5:


(A, B, E, C, B, D)



Paths

A *trivial* path of length 0:

(A)



Simple paths

A simple path is a path in which no vertices are repeated


A simple cycle is a path of at least two vertices with the first and last vertices equal




Connectedness

Two vertices v_i, v_j are said to be **connected** if there exists a path from v_i to v_j

A graph is connected if there exists a path between any two vertices



A connected graph




An unconnected graph

Weighted graphs

A weight may be associated with each edge

- This could represent distance, energy consumption, cost, etc.
- Such a graph is called a *weighted graph*


We will represent weights by numbers next to the edges



Weighted graphs

The *length* of a path within a weighted graph is the sum of all of the weighted edges which make up the path


- The length of the path (A, D, G) in the following graph is $5.1 + 3.7 = 8.8$



Weighted graphs

Different paths may have different length


- Another path is (A, C, F, G) with length $1.2 + 1.4 + 4.5 = 7.1$



Weighted graphs

Problem: Find the shortest path between two vertices


- Here, the shortest path from A to G is (A, C, F, D, E, G) with length 5.7



Trees

A graph is a tree (i) if it is connected and (ii) there is a unique path between any two vertices

- Three trees on the same eight vertices



Consequences:

- The number of edges is $|E| = |V| - 1$
- The graph is *acyclic*, that is, it does not contain any cycles
- Adding one more edge must create a cycle
- Removing any one edge creates two disjoint non-empty sub-graphs

Trees





Any tree can be converted into a rooted tree by:

- Choosing any vertex to be the root
- Defining its neighboring vertices as its children

and then recursively defining:

- All neighboring vertices other than that one designated its parent are now defined to be that vertex's children

Given this tree, here are three rooted trees associated with it



Forests

A forest is an undirected graph, where any two vertices are connected by at most one path


The forest consists of a disjoint union of trees

Consequences:

- The number of edges is $|E| < |V|$
- The number of trees is $|V| - |E|$
- Removing any one edge adds one more tree to the forest

Here is a forest with 22 vertices and 18 edges

- There are four trees



Directed graphs

In **a directed graph**, the edges on a graph are associated with a direction


- Edges are ordered pairs (v_j, v_k) denoting a connection from v_j to v_k
- The edge (v_j, v_k) is different from the edge (v_k, v_j)

Directed graphs

Given our graph of nine vertices $V = \{v_1, v_2, \dots, v_9\}$

- These six pairs (v_j, v_k) are *directed edges*

$$E = \{(v_1, v_2), (v_3, v_5), (v_5, v_3), (v_6, v_9), (v_8, v_4), (v_9, v_4)\}$$



Directed graphs

The maximum number of directed edges in a directed graph is

$$|E| \leq 2 \binom{|V|}{2} = 2 \frac{|V|(|V|-1)}{2} = |V|(|V|-1) = O(|V|^2)$$

In and out degrees


The degree of a vertex in a directed graph:

- The **out-degree** of a vertex is the number of vertices which are adjacent to the given vertex
- The **in-degree** of a vertex is the number of vertices which this vertex is adjacent to

In this graph:

$$\text{in_degree}(v_1) = 0 \quad \text{out_degree}(v_1) = 2$$

$$\text{in_degree}(v_5) = 2 \quad \text{out_degree}(v_5) = 3$$




Sources and sinks

Some definitions:

- Vertices with an in-degree of zero are described as **sources**
- Vertices with an out-degree of zero are described as **sinks**

In this graph:

- Sources: v_1, v_6, v_7
- Sinks: v_2, v_9



Paths

A path in a directed graph is an ordered sequence of vertices

$$(v_0, v_1, v_2, \dots, v_k)$$


where (v_{j-1}, v_j) is an edge for $j = 1, \dots, k$

A path of length 5 in this graph is

$$(v_1, v_4, v_5, v_3, v_5, v_2)$$

A simple cycle of length 3 is


$$(v_8, v_4, v_5, v_8)$$



Weighted directed graphs

In a weighted directed graphs, each edge is associated with a value

Unlike weighted undirected graphs, if both (v_j, v_k) and (v_k, v_j) are edges, the two edges may have different weights




Directed acyclic graphs


A **directed acyclic graph** is a directed graph which has **no cycles**

- These are commonly referred to as **DAG**
- They are graphical representations of partial orders on a finite number of elements

These two are DAG:




This directed graph is not acyclic:



Representations of Adjacency Relations

How do we store the adjacency relations?

- Binary-relation list
- Adjacency matrix
- Adjacency list




Binary-relation list

A **relation list** is the simplest but inefficient:

- A container storing the edges

$\{(1, 2), (1, 4), (3, 5), (4, 2), (4, 5), (5, 2), (5, 3), (5, 8), (6, 9), (7, 9), (8, 4)\}$

- Requires $\Theta(|E|)$ memory
- Determining if v_j is adjacent to v_k is $O(|E|)$
- Finding all neighbors of v_j is $\Theta(|E|)$




Adjacency matrix

An adjacency matrix requires more memory but faster

- The matrix entry (j, k) is set to true if there is an edge (v_j, v_k)

	1	2	3	4	5	6	7	8	9
1		T		T					
2									
3					T				
4		T			T				
5		T	T					T	
6									T
7									T
8				T					
9									

- Requires $\Theta(|V|^2)$ memory
- Determining if v_j is adjacent to v_k is $O(1)$
- Finding all neighbors of v_j is $\Theta(|V|)$



Adjacency list

An adjacency list:

- Each vertex is associated with a list of its neighbors


```

1  • → 2 → 4
2  •
3  • → 5
4  • → 2 → 5
5  • → 2 → 3 → 8
6  • → 9
7  • → 9
8  • → 4
9  •
  
```

- Requires $\Theta(|V| + |E|)$ memory

- On average:

- Determining if v_j is adjacent to v_k is $O\left(\frac{|E|}{|V|}\right)$
- Finding all neighbors of v_j is $\Theta\left(\frac{|E|}{|V|}\right)$



The Graph ADT

The Graph ADT describes a container storing an adjacency relation

- Queries include:
 - The number of vertices
 - The number of edges
 - List the vertices adjacent to a given vertex
 - Are two vertices adjacent?
 - Are two vertices connected?
- Modifications include:
 - Inserting or removing an edge
 - Inserting or removing a vertex (and all edges containing that vertex)

The runtime of these operations will depend on the representation

Summary

In this topic, we have covered:

- Basic graph definitions
 - Vertex, edge, degree, adjacency
- Paths, simple paths, and cycles
- Connectedness
- Weighted graphs
- Directed graphs
- Directed acyclic graphs

We will continue by looking at a number of problems related to graphs

References

Wikipedia, http://en.wikipedia.org/wiki/Adjacency_matrix
http://en.wikipedia.org/wiki/Adjacency_list

- [1] Donald E. Knuth, *The Art of Computer Programming, Volume 1: Fundamental Algorithms*, 3rd Ed., Addison Wesley, 1997, §2.2.1, p.238.
- [2] Cormen, Leiserson, and Rivest, *Introduction to Algorithms*, McGraw Hill, 1990, §11.1, p.200.
- [3] Weiss, *Data Structures and Algorithm Analysis in C++*, 3rd Ed., Addison Wesley, §3.6, p.94.
- [4] David H. Laidlaw, Course Notes, <http://cs.brown.edu/courses/cs016/lectures/13%20Graphs.pdf>